

Review Diffusion model

Problem Formulation: To generate image, we gradually add noise to a given set of images in order to transform it a white noise (Gaussian distribution with mean 0 and variance \mathbf{I}).

Forward process: Adding noise to given sample x_0 gradually and expecting that at the end of process $x_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.

$$q(x_{1:T}|x_0) = q(x_0) \cdot \prod_1^T q(x_t|x_{t-1})$$

$$q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} x_{t-1}, \beta_t \mathbf{I})$$

Set $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$, we have:

$$q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{\alpha_t} x_{t-1}, (1 - \alpha_t) \mathbf{I})$$

The special point is that, using that notation gives the ability of sampling x_t at arbitrary timestep t , which is achieved by transformation below:

$$x_t = \sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} z_{t-1} \quad (1)$$

$$= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_{t-1}} z_{t-2}) + \sqrt{1 - \alpha_t} z_{t-1} \quad (2)$$

$$= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \sqrt{\alpha_t (1 - \alpha_{t-1})} z_{t-2} + \sqrt{1 - \alpha_t} z_{t-1} \quad (3)$$

Where $z_{t-1}, z_{t-2} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. Think $\sqrt{\alpha_t (1 - \alpha_{t-1})} z_{t-2}$ as random variable u and $\sqrt{1 - \alpha_t} z_{t-1}$ as random variable v , we have $u \sim \mathcal{N}(\mathbf{0}, \alpha_t (1 - \alpha_{t-1}) \mathbf{I})$ and $v \sim \mathcal{N}(\mathbf{0}, (1 - \alpha_t) \mathbf{I})$. Since u and v are independent, $V[u+v] = V[u] + V[v]$ and we have $(u + v) \sim \mathcal{N}(\mathbf{0}, (1 - \alpha_t \alpha_{t-1}) \mathbf{I})$. Then equation (3) can be rewrited as follow:

$$\mathbf{x}_t = \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + (u + v) \quad (4)$$

$$= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} z \quad (5)$$

Where $z \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, so we have $x_t \sim \mathcal{N}(\sqrt{\alpha_t \alpha_{t-1}} x_{t-2}, (1 - \alpha_t \alpha_{t-1}) \mathbf{I})$. Do this transformation repeatedly and we have closed form for sampling x_t from x_0 as follow:

$$\mathbf{x}_t = \sqrt{\alpha_t \alpha_{t-1} \dots \alpha_1} x_0 + \sqrt{1 - \alpha_t \alpha_{t-1} \dots \alpha_1} z \quad (6)$$

$$= \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} z \quad (7)$$

Where $z \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, so $x_t \sim \mathcal{N}(\sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) \mathbf{I})$.

Background for deriving loss function:

What is ELBO (Evidence Lower Bound)?

Let say that in a latent variable model, we posit that our observed data x is a realization of a random variable X . Moreover, we posit the existence of another random variable Z where X and Z are distributed according to a joint distribution $p(X, Z, \theta)$. The Z variable does not have any observation so remains a latent variable. With that model, there are two predominate tasks which are interesting to dig in:

- Given some fixed value for θ , compute the posterior distribution $p(Z|X, \theta)$.
- Given θ is unknown, find the maximum likelihood estimate of θ : $\text{argmax}_{\theta} l(\theta)$, where $l(\theta)$ is the log likelihood function:

$$\begin{aligned} l(\theta) &= \log(p(x, \theta)) \\ &= \log \int_z p(x, z, \theta) dz \end{aligned}$$

The term "evidence" is just a name given to the likelihood function $l(\theta) = \log(x, \theta)$, so "evidence lower bound" is the lower bound of likelihood function, which can be derived as follow:

$$\begin{aligned} \log(x, \theta) &= \log \int_z p(x, z, \theta) dz \\ &= \log \int_z q(z) \frac{p(x, z, \theta)}{q(z)} dz \\ &= \log \left(\mathbf{E}_{z \sim q} \left[\frac{p(x, z, \theta)}{q(z)} \right] \right) \\ &\geq \mathbf{E}_{z \sim q} \left[\log \left(\frac{p(x, z, \theta)}{q(z)} \right) \right] \end{aligned}$$

The last inequality follows from Jensen inequality. The gap between evidence and ELBO is KL divergence:

$$\begin{aligned} \log(p(x, \theta)) - \mathbf{E}_{z \sim q} \left[\log \left(\frac{p(x, z, \theta)}{q(z)} \right) \right] &= \mathbf{E}_{z \sim q} [\log(p(x, \theta))] - \mathbf{E}_{z \sim q} \left[\log \left(\frac{p(x, z, \theta)}{q(z)} \right) \right] \\ &= \mathbf{E}_{z \sim q} \left[\log(p(x, \theta)) - \log \left(\frac{p(x, z, \theta)}{q(z)} \right) \right] \\ &= \mathbf{E}_{z \sim q} \left[\log \left(p(x, \theta) \frac{q(z)}{p(x, z, \theta)} \right) \right] \\ &= \mathbf{E}_{z \sim q} \left[\log \frac{q(z)}{p(z|x, \theta)} \right] \\ &= \mathbf{D}_{\text{KL}}(q(z) || p(z|x, \theta)) \end{aligned}$$

Training is procedure of maximizing negative log likelihood, which can be achieved approximately by maximizing ELBO (variational bound).

$$\mathbf{E}[-\log p_\theta(x_0)] \leq \mathbf{E}_q \left[-\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right]$$

Set $L = \mathbf{E}_q \left[-\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right]$ and make a few transformation, we have:

$$\begin{aligned} L &= \mathbf{E}_q \left[-\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right] \\ &= \mathbf{E}_q \left[-\log p(x_T) - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] \\ &= \mathbf{E}_q \left[-\log p(x_T) - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} - \log \left(\frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right) \right] \\ &= \mathbf{E}_q \left[-\log p(x_T) - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \cdot \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} - \log \left(\frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right) \right] \text{ (apply bayes rule)} \\ &= \mathbf{E}_q \left[-\log p(x_T) - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \sum_{t \geq 1} \log \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} - \log \left(\frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right) \right] \\ &= \mathbf{E}_q \left[-\log p(x_T) - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \log \frac{1}{q(x_T|x_0)} - \log \left(\frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right) \right] \\ &= \mathbf{E}_q \left[-\log \frac{p(x_T)}{q(x_T|x_0)} - \sum_{t \geq 1} \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \log \left(\frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} \right) \right] \end{aligned}$$

Let $L_t = \mathbf{E}_q \left[-\log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \right]$ as an example.

$$\begin{aligned} \mathbf{E}_q \left[-\log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \right] &= \int q(x_{0:T}) \log \frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} dx_{0:T} \\ &= \int q(x_{t-1}|x_t, x_0) q(x_t, x_0) q(x_{1:t-2}, x_{t+1:T}|x_0, x_{t-1}, x_t) \log \frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} dx_{0:T} \\ &= \int \left[\int q(x_{t-1}|x_t, x_0) \log \frac{q(x_{t-1}|x_t, x_0)}{p_\theta(x_{t-1}|x_t)} dx_{t-1} \right] q(x_{0:t-2}, x_{t:T}|x_{t-1}) dx_{0:t-2, t:T} \\ &= \int [\text{D}_{\text{KL}}(q(x_{t-1}|x_0, x_t) || p_\theta(x_{t-1}|x_t))] q(x_{0:t-2}, x_{t:T}|x_{t-1}) dx_{0:t-2, t:T} \\ &= \mathbf{E}_{q(x_{0:t-2}, x_{t:T}|x_{t-1})} [\text{D}_{\text{KL}}(q(x_{t-1}|x_0, x_t) || p_\theta(x_{t-1}|x_t))] \end{aligned}$$

One noticeable property is that with given $x_{t-1}, x_0, q(x_{t-1}|x_0, x_t)$ and $p_\theta(x_{t-1}|x_t)$ are completely defined and $D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))$ is known, or can be set as a constant. Thus, we have the following equation:

$$\mathbf{E}_{q(x_{t-1})}[D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))] = D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))$$

From that observation, we finally get the form of L_t :

$$\begin{aligned} \mathbf{E}_q \left[-\log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \right] &= \mathbf{E}_{q(x_0:t-2, x_{t:T}|x_{t-1})}[D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))] \\ &= \mathbf{E}_{q(x_0:t-2, x_{t:T}|x_{t-1})} \mathbf{E}_{q(x_{t-1})}[D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))] \\ &= \mathbf{E}_{q(x_0:T)}[D_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t))] \end{aligned}$$

Applying above transformation to the general L , we get the final form of L :

$$\begin{aligned} L &= \mathbf{E}_q \left[-\log \frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right] \\ &= \mathbf{E}_q \left[D_{\text{KL}}(q(x_T|x_0)||p(x_T)) + \sum_{t>1} D_{\text{KL}}(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) + \log\left(\frac{q(x_1|x_0)}{p_\theta(x_0|x_1)}\right) \right] \end{aligned}$$

We can formulate form of posterior $q(x_{t-1}|x_t, x_0)$ base on Bayes rule.

$$q(x_{t-1}|x_t, x_0) = \frac{q(x_t, x_{t-1}|x_0)}{q(x_t|x_0)} \quad (8)$$

$$= q(x_t|x_{t-1}, x_0) \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} \quad (9)$$

$$= q(x_t|x_{t-1}) \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} \quad (10)$$

$$\propto \exp \left(-\frac{1}{2} \left(\frac{(x_t - \sqrt{\alpha_t} x_{t-1})^2}{\beta_t} + \frac{(x_{t-1} - \sqrt{\bar{\alpha}_{t-1}} x_0)^2}{1 - \bar{\alpha}_{t-1}} - \frac{(x_t - \sqrt{\bar{\alpha}_t} x_0)^2}{1 - \bar{\alpha}_t} \right) \right) \quad (11)$$

(Shout out for Cuong Pham)

$$= \exp \left(-\frac{1}{2} \left(\left(\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}} \right) x_{t-1}^2 - \left(\frac{2\sqrt{\alpha_t}}{\beta_t} x_t + \frac{2\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} x_0 \right) x_{t-1} + \mathbf{C}(x_t, x_0) \right) \right) \quad (12)$$

Where $C(x_t, x_0)$ is a term not relating to x_t . Equation ?? can be derived to Gaussian

form. Making a simple transformation to do that.

$$\left(\frac{\alpha_t}{\beta_t} + \frac{1}{1 - \bar{\alpha}_{t-1}}\right)x_{t-1}^2 - \left(\frac{2\sqrt{\alpha_t}}{\beta_t}x_t + \frac{2\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}}x_0\right)x_{t-1} + \mathbf{C}(x_t, x_0) \quad (13)$$

$$= \left(\frac{\alpha_t(1 - \bar{\alpha}_{t-1}) + \beta_t}{\beta_t(1 - \bar{\alpha}_{t-1})}\right)x_{t-1}^2 - 2\left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})x_t + \sqrt{\bar{\alpha}_{t-1}}\beta_t x_0}{\beta_t(1 - \bar{\alpha}_{t-1})}\right)x_{t-1} + \mathbf{C}(x_t, x_0) \quad (14)$$

$$= \frac{\alpha_t - \alpha_t\bar{\alpha}_{t-1} + 1 - \alpha_t}{\beta_t(1 - \bar{\alpha}_{t-1})}x_{t-1}^2 - 2\left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{\beta_t(1 - \bar{\alpha}_{t-1})}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\beta_t(1 - \bar{\alpha}_{t-1})}x_0\right)x_{t-1} + \mathbf{C}(x_t, x_0) \quad (15)$$

$$= \frac{1 - \bar{\alpha}_t}{\beta_t(1 - \bar{\alpha}_{t-1})}x_{t-1}^2 - 2\left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{\beta_t(1 - \bar{\alpha}_{t-1})}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{\beta_t(1 - \bar{\alpha}_{t-1})}x_0\right)x_{t-1} + \mathbf{C}(x_t, x_0) \quad (16)$$

$$= \frac{1 - \bar{\alpha}_t}{\beta_t(1 - \bar{\alpha}_{t-1})} \left(x_{t-1}^2 - 2\left(\frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0\right)x_{t-1} \right) + \mathbf{C}(x_t, x_0) \quad (17)$$

From here, we have $\mathbf{E}_{q(x_{t-1}|x_t, x_0)}[x_{t-1}] = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0$ and $\mathbf{V}_{q(x_{t-1}|x_t, x_0)}[x_{t-1}] = \frac{\beta_t(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}$.

Training now is to find $p_\theta(x_{t-1}|x_t)$ that approximates $q(x_{t-1}|x_t, x_0)$. It is straight to form the reverse probability as a Gaussian distribution: $\mathcal{N}(x_{t-1}, \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$.

According to DDPM, the term Σ_θ can be fixed and the model can even create high-quality samples. In detail, it is fixed to $\sigma_t^2 \mathbf{I} = \tilde{\beta} \mathbf{I}$. The critical part is to define an effective way to approximate the μ_θ . As the posterior and the approximate reverse distribution has the same form, we can rewrite the loss as follow.

$$\begin{aligned} \text{D}_{\text{KL}}(q(x_{t-1}|x_0, x_t)||p_\theta(x_{t-1}|x_t)) &= \frac{1}{2} \log \frac{\sigma_t}{\tilde{\sigma}_t} + \frac{\tilde{\sigma}_t^2 + \|\tilde{\mu} - \mu_\theta\|^2}{2\sigma_t^2} - \frac{1}{2} \\ &= \frac{\|\tilde{\mu} - \mu_\theta\|^2}{2\sigma_t^2} + C \end{aligned}$$

(Due to the assumption of $\sigma_t^2 \mathbf{I} = \tilde{\beta} \mathbf{I}$)